

Exploratory data analysis, EDA

Nazgul Rakhimzhanova



COURSE SCHEDULE

| week | Mid Term (weeks 01-07) | End Term (weeks 08-14) | week |
|-----------|---|--------------------------------|-----------|
| 01 | Intro: Data Science Area and open source tools for Data Science | | 08 |
| 02 | NumPy package for data science | Sampling and Estimation | 09 |
| 03 | Pandas package for data science | Correlation and Covariance | 10 |
| 04 | Visualization with matplotlib | Hypothesis testing | 11 |
| 05 | Statistics: Distribution – Normal | Decision Tree | 12 |
| 06 | Exploratory Data Analysis (EDA) | Linear Regression | 13 |
| <u>07</u> | Summary for 6 weeks QA session | Summary for 6 weeks QA session | <u>14</u> |
| 15 | Course s | ummary | |

OUTLINE



- Previously
- □ EDA
- Readings

PREVIUOSLY



- Discussed about distributions
- Did overview of the EDA
- Practiced with histograms to identify he distribution



PREVIUOSLY

Have you thought about additional DS application for real life tasks?



What We Do

- Data Collection
- EDA
- Feature Engineering
- Model Training
- Model Improvement

EDA common steps



- Data specification (understanding)
 - First look at data
 - How many data types?
 - How many missing values?
- Handling missing values
 - Drop / restore
- Data editing (correction)
 - Inconvenient formats
 - Messy data noise, outliers.
 - Categorical data handling
- Relationships
 - Patterns, pre-summaries
- Normalization
- Feature Extraction

DATA SPECIFICATION



- •Understanding the data
 - pandas.dataframe.head()
 - •pandas.dataframe.tail()
 - pandas.dataframe.info()
 - •pandas.dataframe.describe()
 - •pandas.dataframe['column'].value_counts()

DATA UNDERSTANDING



- •Understanding -> Visualization
 - •pandas.dataframe.hist()
 - •pandas.boxplot()
 - •matplotlib etc.



- Formats -> Date to DateTime
- <u>Dropping inconsistent</u> data -> Dropping IDs and similar data with no knowledge.
- Dropping NaNs: rule >5% -- 30%<
 - df.drop('column', axis=1, inplace = True)
 - df.drop(np.arrange(10), axis=0, inplace = True)
 - •df.drop(df[(df['age'] < 18) | (df['age'] >
 50)].index, axis = 0, inplace = True)

HANDLING MISSING VALUES



- •Missing values of column
- •Missing values of rows

• How to restore?



- Restoring NaN values:
 - Restore by mean
 - Restore by category mean
 - Restore by median
 - Restore by mode
 - Restore by sliding windows' mean/median/mode
 - Restore by backward/forward replication
 - Restore by interpolation
 - Restore by filling with 0
 - NaN != 0





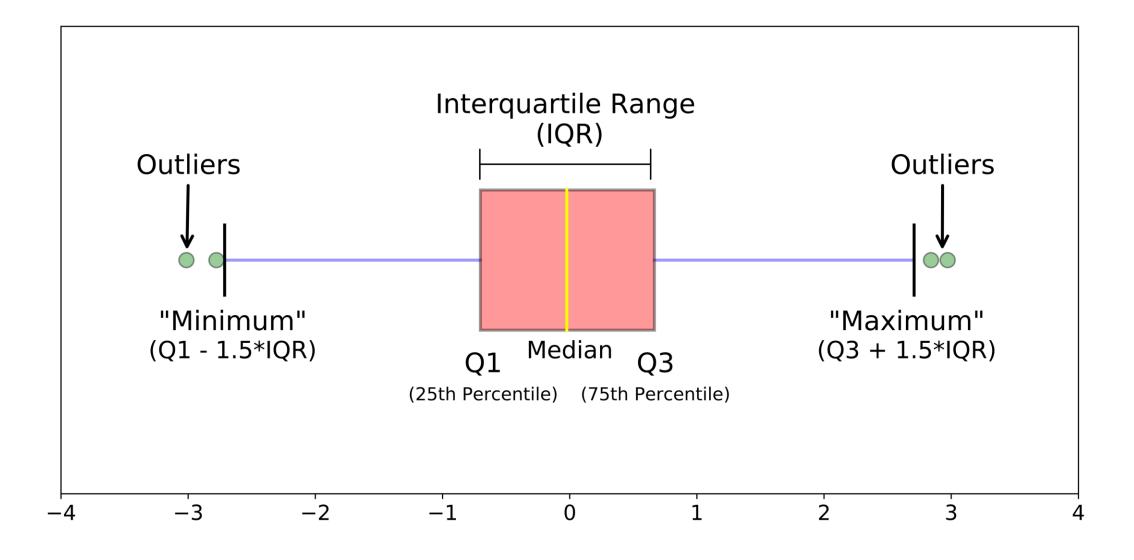
OUTLIERS



- -Boxplots
- -Histograms
- -IQR

OUTLIERS





CATEGORICAL DATA ENCODING

- Label Encoding
- •One-Hot Encoding
- Hashing (HASH function)



Label Encoding

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | duration | caı |
|---|-----|-------------|---------|-------------------|---------|---------|------|-----------|-------|-------------|----------|-----|
| 0 | 26 | student | single | high.school | no | no | no | telephone | jun | mon | 901 | 1 |
| 1 | 46 | admin. | married | university.degree | no | yes | no | cellular | aug | tue | 208 | 2 |
| 2 | 49 | blue-collar | married | basic.4y | unknown | yes | yes | telephone | jun | tue | 131 | 5 |
| 3 | 31 | technician | married | university.degree | no | no | no | cellular | jul | tue | 404 | 1 |
| 4 | 42 | housemaid | married | university.degree | no | yes | no | telephone | nov | mon | 85 | 1 |

```
education
university - 6
professional courses - 5
college - 4
high school - 3
basic 12y - 2
basic 9y - 1
basic 4y - 0
```

```
martial
Single - 2
Married - 1
```

$$\frac{\text{loan}}{\text{no} - 0}$$

$$\text{yes} - 2$$



Label Encoding

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | duration | caı |
|---|-----|-------------|---------|-------------------|---------|---------|------|-----------|-------|-------------|----------|-----|
| 0 | 26 | student | single | high.school | no | no | no | telephone | jun | mon | 901 | 1 |
| 1 | 46 | admin. | married | university.degree | no | yes | no | cellular | aug | tue | 208 | 2 |
| 2 | 49 | blue-collar | married | basic.4y | unknown | yes | yes | telephone | jun | tue | 131 | 5 |
| 3 | 31 | technician | married | university.degree | no | no | no | cellular | jul | tue | 404 | 1 |
| 4 | 42 | housemaid | married | university.degree | no | yes | no | telephone | nov | mon | 85 | 1 |

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | duration | campaign |
|---|-----|-----|---------|-----------|---------|---------|------|---------|-------|-------------|----------|----------|
| 0 | 26 | 8 | 2 | 3 | 0 | 0 | 0 | 1 | 4 | 1 | 901 | 1 |
| 1 | 46 | 0 | 1 | 6 | 0 | 2 | 0 | 0 | 1 | 3 | 208 | 2 |
| 2 | 49 | 1 | 1 | 0 | 1 | 2 | 2 | 1 | 4 | 3 | 131 | 5 |
| 3 | 31 | 9 | 1 | 6 | 0 | 0 | 0 | 0 | 3 | 3 | 404 | 1 |
| 4 | 42 | 3 | 1 | 6 | 0 | 2 | 0 | 1 | 7 | 1 | 85 | 1 |



•One-Hot Encoding

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | duration | caı |
|---|-----|-------------|---------|-------------------|---------|---------|------|-----------|-------|-------------|----------|-----|
| 0 | 26 | student | single | high.school | no | no | no | telephone | jun | mon | 901 | 1 |
| 1 | 46 | admin. | married | university.degree | no | yes | no | cellular | aug | tue | 208 | 2 |
| 2 | 49 | blue-collar | married | basic.4y | unknown | yes | yes | telephone | jun | tue | 131 | 5 |
| 3 | 31 | technician | married | university.degree | no | no | no | cellular | jul | tue | 404 | 1 |
| 4 | 42 | housemaid | married | university.degree | no | yes | no | telephone | nov | mon | 85 | 1 |

| Single | Married |
|--------|---------|
| 1 | 0 |
| 0 | 1 |
| 0 | 1 |
| 0 | 1 |
| 0 | 1 |
| 0 | 1 |
| 1 | 0 |



•One-Hot Encoding

| | age | job | marital | education | default | housing | loan | contact | month | day_of_week | duration | caı |
|---|-----|-------------|---------|-------------------|---------|---------|------|-----------|-------|-------------|----------|-----|
| 0 | 26 | student | single | high.school | no | no | no | telephone | jun | mon | 901 | 1 |
| 1 | 46 | admin. | married | university.degree | no | yes | no | cellular | aug | tue | 208 | 2 |
| 2 | 49 | blue-collar | married | basic.4y | unknown | yes | yes | telephone | jun | tue | 131 | 5 |
| 3 | 31 | technician | married | university.degree | no | no | no | cellular | jul | tue | 404 | 1 |
| 4 | 42 | housemaid | married | university.degree | no | yes | no | telephone | nov | mon | 85 | 1 |

| | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 43 | 44 | 45 | 46 | 47 | 48 | 49 | 50 | 51 | 52 |
|---|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|---------|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0.0 | 1.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 1 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 2 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 3 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 1.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |
| 4 | 0.0 | 1.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 1.0 | 0.0 |



- -Patters
 - -Barplots to compare the same parameter in two or more groups
 - -GroupBy to see some specific groups or parameters
- -Correlations
 - -corr() to check if our data happen to have linear dependency



"Never trust a statistic if you haven't falsified it yourself".

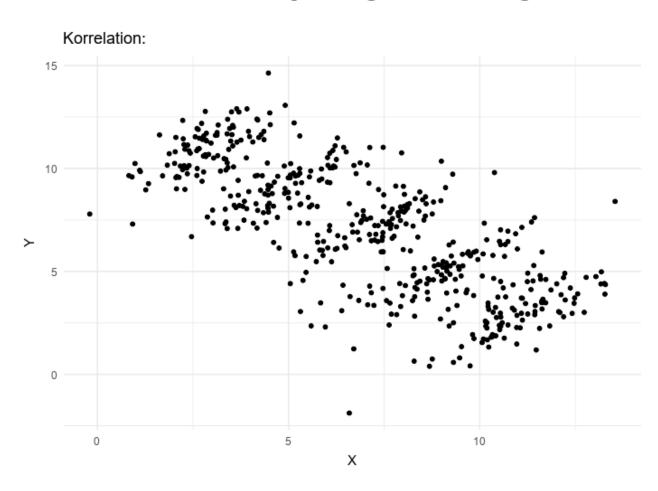
Winston Churchill

The Simpson-Paradox

Simpson's paradox, which also goes by several other names, is a phenomenon in probability and statistics, in which a trend appears in several different groups of data but disappears or reverses when these groups are combined.

https://en.wikipedia.org/wiki/Simpson%27s_paradox





https://en.wikipedia.org/wiki/Simpson%27s_paradox





Fact 01:

| Gender | Applicants | Accepted | Rejected | % Accepted |
|--------|------------|----------|----------|------------|
| Male | 2175 | 1025 | 1150 | 0.471 |
| Female | 849 | 261 | 588 | 0.307 |

University of California, Berkeley, Admission rate 1973 y.



Fact 02:

| Gender | Applicants | Accepted | Rejected | % Accepted |
|--------|------------|----------|----------|------------|
| Male | 825 | 512 | 313 | 0.621 |
| Female | 108 | 89 | 19 | 0.824 |

| Gender | Applicants | Accepted | Rejected | % Accepted |
|--------|------------|----------|----------|------------|
| Male | 417 | 138 | 279 | 0.331 |
| Female | 375 | 131 | 244 | 0.349 |

Data for Department 1.

Data for Department 3.

| Gender | Applicants | Accepted | Rejected | % Accepted |
|--------|------------|----------|----------|------------|
| Male | 560 | 353 | 207 | 0.63 |
| Female | 25 | 17 | 8 | 0.68 |

| Gender | Applicants | Accepted | Rejected | % Accepted |
|--------|------------|----------|----------|------------|
| Male | 373 | 22 | 351 | 0.059 |
| Female | 341 | 24 | 317 | 0.070 |

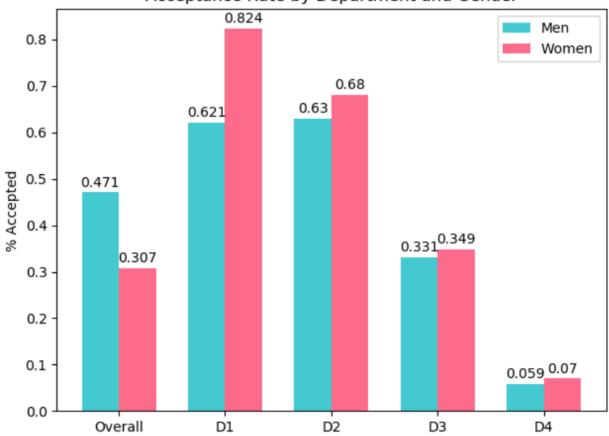
Data for Department 2.

Data for Department 4.

The overall acceptance rate of women is lower than the overall acceptance rate of men. Yet, in each department, the acceptance rate for women is higher than the acceptance rate for men.

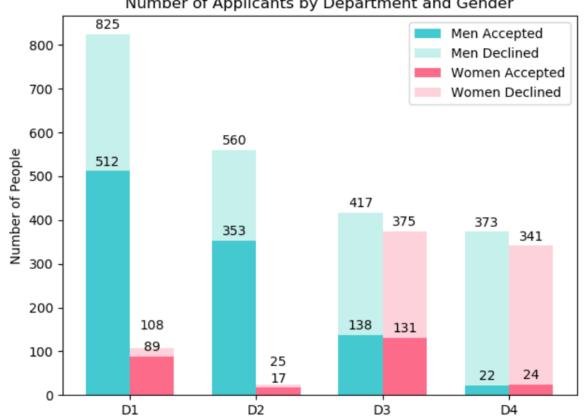












- Department 1 seemed to accept both a high number of people and a high percentage of applicants, yet very few women applied.
- The same goes for Department 2.
- In Departments 3 and 4, the number of women who applied was almost the same as the number of men who applied — but the overall acceptance rate was quite low compared to the other departments.

This is also the explanation of why the overall acceptance rate for women is lower than the rate for men. It's not that women were discriminated against by any departments (at least as far as we know!), it's that women — in comparison to men — seemed to apply more to very competitive departments where it was hard to get in.

Readings



- https://towardsdatascience.com/exploratory-data-analysis-of-kaggle-datasets-9a293886f644
- https://towardsdatascience.com/gender-bias-in-admission-statistics-the-simpson-paradox-cd381d994b16